

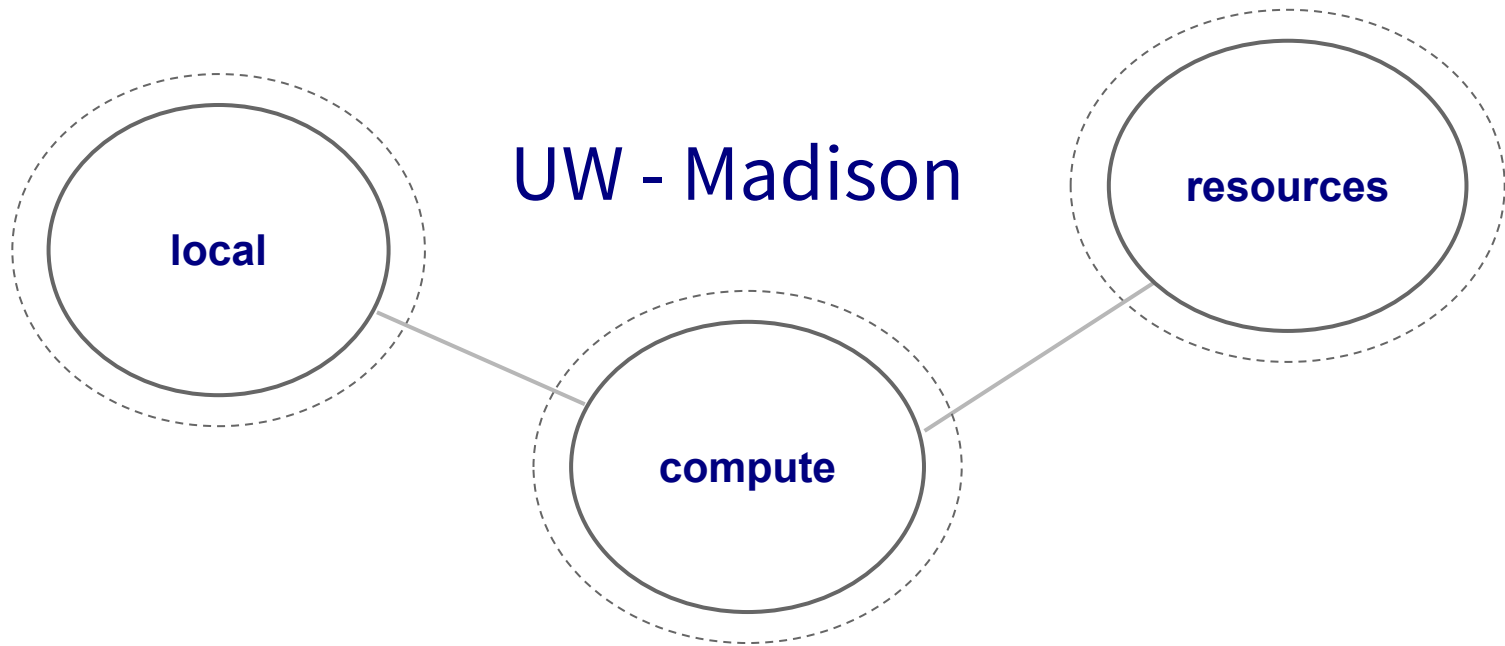
Introduction to DHTC

Brian Lin


OSG Software Team

University of Wisconsin - Madison

Local High Throughput Computing



Local High Throughput Computing





How do you get more computing resources?



#1: Buy Hardware

#1: Buy Hardware

- Great for specific hardware/privacy requirements
- Costs \$\$\$
 - Initial cost
 - Maintenance
 - Management
 - Power and cooling
- Delivery and installation takes time
- Rack/floor space
- Obsolescence
- Plan for peak usage, pay for all usage



#2: Use the Cloud

#2: Use the Cloud - Pay per cycle

- e.g. Amazon Web Services, Google Compute Engine, Microsoft Azure, Rackspace
- Fast spin-up
- Costs \$\$\$
- Still needs expertise + management
 - Easier than in the past with the `condor_annex` tool
- Does payment fit with your institutional or grant policies?


#2: Use the Cloud - ‘Managed’ clouds

- e.g. Cycle Computing, Globus Genomics
- Pay someone to manage your cloud resources — still costs \$\$\$
- Researchers and industry have used this to great success
 - [Using Docker, HTCondor, and AWS for EDA Model Development](#)
 - [Optimizations in running large-scale Genomics workloads in Globus Genomics using HTCondor](#)
 - [HTCondor in the enterprise](#)
 - [HTCondor at Cycle Computing: Better Answers. Faster.](#)



#3: Share Resources

#3: Share Resources - Distributed HTC





i.

Split Up Your Jobs Manually


Let's start sharing!

Manual Job Split




- Obtain login access
- Query each cluster for idle resources
- Split and submit jobs based on resource availability

#3: Share Resources - Distributed HTC



#3: Share Resources - Distributed HTC



Manual Job Split - Shortcomings

- Fewer logins = fewer potential resources
- More logins = more account management
- Why would they give you accounts? Are your friends going to want CHTC accounts?
- Querying and splitting jobs is tedious and inaccurate
- Not all clusters use HTCondor — other job schedulers e.g., SLURM, PBS, etc.
- Pools are independent — workflows must be confined to a single pool



ii.

Split Up Your Jobs Automatically

Let the computers do the work

Automatic Job Split - Shortcomings



Homer: Kids: there's three ways to do things; the right way, the wrong way and the Max Power way!


Bart: Isn't that the wrong way?

Homer: Yeah, but faster!

Groening, M (Writer), Michels, P. (Director) . (1999).
Homer to the Max [Television Series Episode]. In
Scully, M. (Executive Producer), *The Simpsons*. Los
Angeles, CA: Gracie Films

Automatic Partitions - Shortcomings

"I SPEND A LOT OF TIME ON THIS TASK.
I SHOULD WRITE A PROGRAM AUTOMATING IT!"



#3: Share Resources - Requirements

- Minimal account management
- No job splitting
- DAG workflow functionality
- HTCondor only!
- No resource sharing requirements

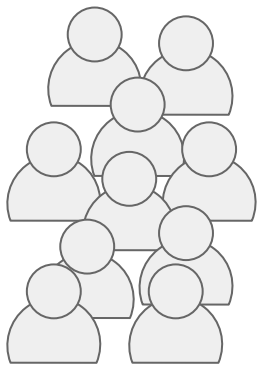


iii.



Overlay Systems

Let the OSG do the heavy lifting

The OSG Model




OSG Submit
and CM




Cluster


The OSG Model



The OSG Model




The OSG Model




Job Matching

- On a regular basis, the central manager reviews Job and Machine attributes and matches jobs to slots.



The OSG Model



The OSG Model - Jobs in Jobs



Photo Credit: Shereen M, Untitled, Flickr <https://www.flickr.com/photos/shereen84/2511071028/> (CC BY-NC-ND 2.0)

#3: Share Resources - Requirements

- Minimal account management: only one submit server
- No job splitting: only one HTCondor pool
- DAG workflow functionality: Only one HTCondor pool
- HTCondor only: Only one HTCondor pool
- No resource sharing requirements: the OSG doesn't require that users “pay into” the OSG

The OSG Model - Recap

- Pilot jobs (or pilots) are special jobs
- Pilots are sent to sites with idle resources
- Pilot payload = HTCondor execute node software
- Pilot execute node reports to your OSG pool
- Pilots lease resources:
 - Lease expires after a set amount of time or lack of demand
 - Leases can be revoked!

The OSG Model - Leasing the Cloud

- What if there aren't enough idle resources?
- Combine overlay system with cloud technology
- Some of your OSG jobs may run in the cloud in the next few years
- ... but this should be completely transparent to you

The OSG Model - Collection of Pools

- Your OSG pool is just one of many
- Separate pools for each Virtual Organization (VO)
- Your jobs will run on the OSG VO pool



The OSG Model - Getting Access

- During the school:
 - OSG submit node at UW (exercises)
 - OSG submit node via OSG Connect (exercises)
- After the school:
 - Both of the above
 - Institution-hosted submit node
 - VO-hosted submit nodes



Questions?



Overlay Systems are Awesome!



What's the Catch?

Requires more infrastructure, software, set-up,
management, troubleshooting...

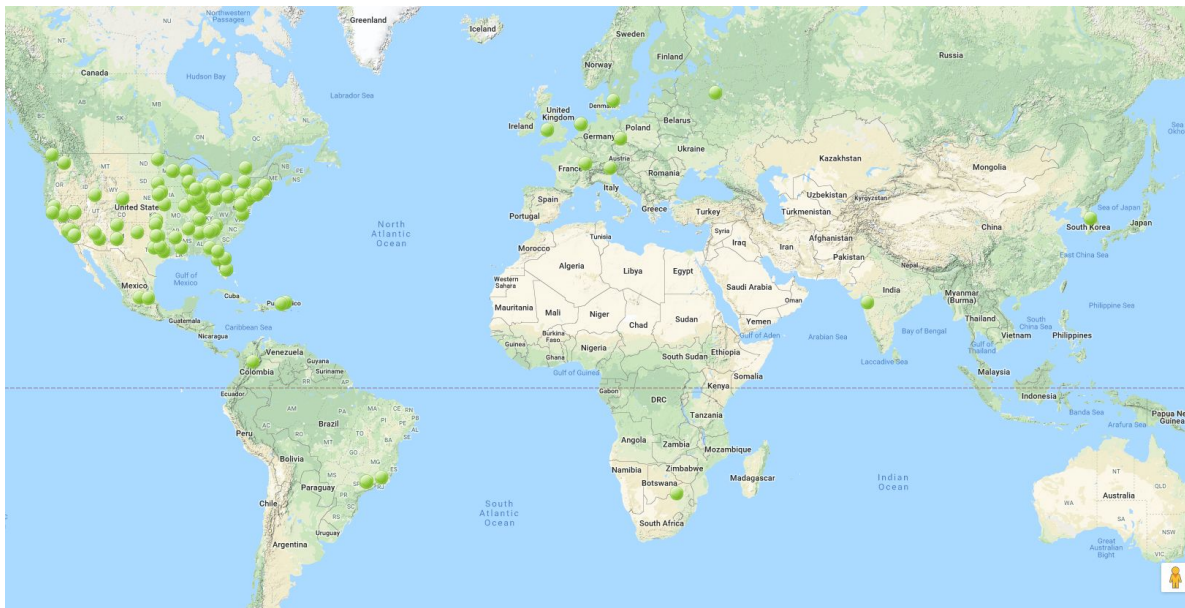
*“You know you have a **distributed system** when the crash of a computer you’ve never heard of stops you from getting any work done.”*

- Leslie Lamport

#1: Heterogenous Resources

Accounting for differences between the
OSG and your local cluster

Sites of the OSG



Source: <http://display.opensciencegrid.org/>

Heterogeneous Resources - Software

- Different operating systems (Red Hat, CentOS, Scientific Linux; versions 6 and 7)
- Varying software versions (e.g., at least Python 2.6)
- Varying software availability (e.g., no BLAST*)

Solution: Make your jobs more portable: OASIS, containers, etc (more in Wednesday's talks)

Hetero. Resources - Hardware

- CPU: Mostly single core
- RAM: Mostly < 8GB
- GPU: Limited #s but more being added
- Disk: No shared file system (more in Thursday's talks)

Solution: Split up your workflow to make your jobs more high throughput

#2: With Great Power Comes Great Responsibility

How to be a good netizen

Resources You Don't Own

- Primary resource owners can kick you off for any reason
- No local system administrator relationships
- No sensitive data (again)!



Be a Good Netizen!

- Use of shared resources is a privilege
- Only use the resources that you request
- Be nice to your submit nodes

Solution: Test jobs on local resources with
`condor_submit -i`



#3: Slower Ramp Up

Leasing resources takes time!

Slower Ramp Up

- Adding slots: pilot process in the OSG vs slots already in your local pool
- Not a lot of time (~minutes) compared to most job runtimes (~hours)
 - Small trade-off for increased availability
 - Tip: If your jobs only run for $< 10\text{min}$ each, consider combining them so each job runs for at least 30min



Robustify Your Jobs

Succeeding in the face of failure

Job Robustification

- Test small, test often
- Specify output, error, and log files at least while you develop your workflow
- Use `on_exit_hold` to catch different failure modes
 - `on_exit_hold = (ExitCode != 3)`
 - `on_exit_hold = (time() - JobCurrentStartDate < 1 * $(HOUR))`
- For jobs that run too long:

```
periodic_hold      = (time() - JobCurrentStartDate > 4 * $(HOUR))
periodic_release   = (HoldReasonCode == 3) && (NumJobStarts < 3)
```

HoldReasonCode is 3 for any jobs where `on_exit_hold` or `periodic_hold` evaluate to True

Job Robustification

- In your own code:
 - Self checkpointing
 - Different exit codes for use with `on_exit_hold`
 - Defensive troubleshooting (`hostname`, `ls -l`, `pwd`, `condor_version` in your wrapper script)
 - Add simple logging (e.g. `print`, `echo`, etc)



Questions?